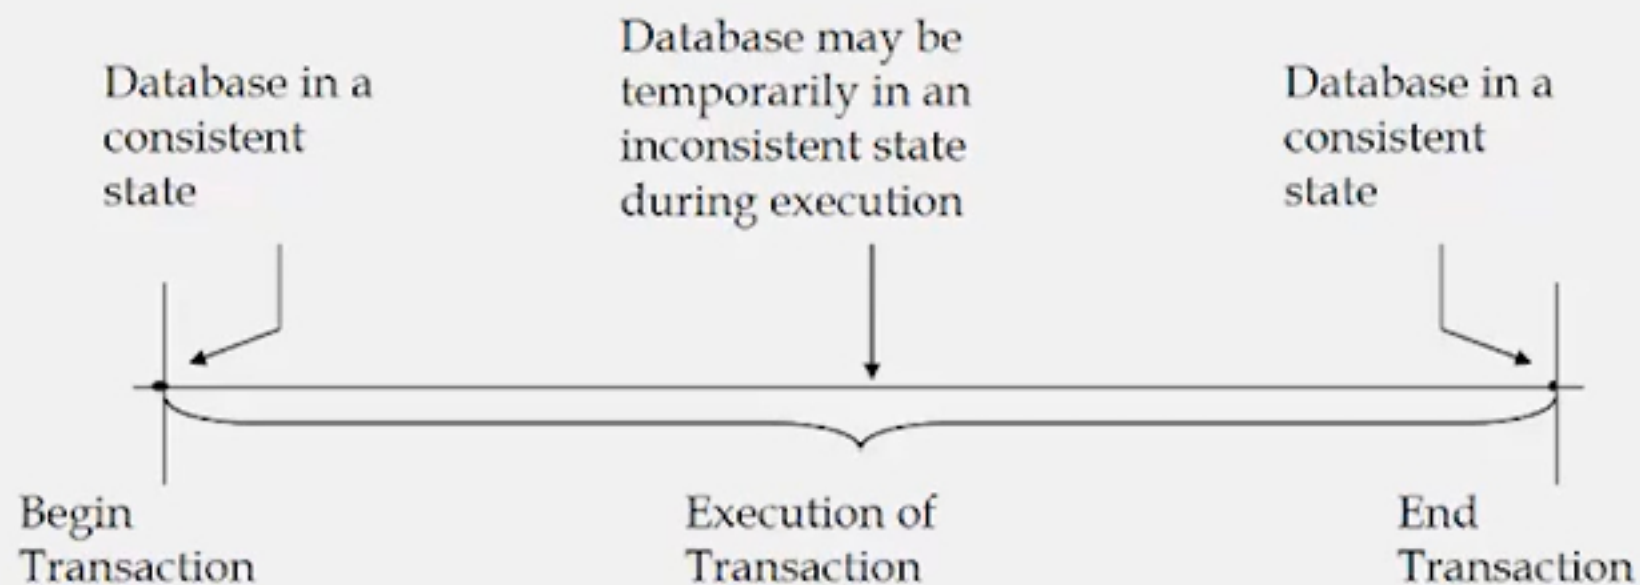




Introduction to Database System Recovery

Transaction

A transaction is a collection of actions that make consistent transformations of system states while preserving system consistency.



Principles of Transactions

ATCOMICITY

- all or nothing

CONSISTENCY

- no violation of integrity constraints

ISOLATION

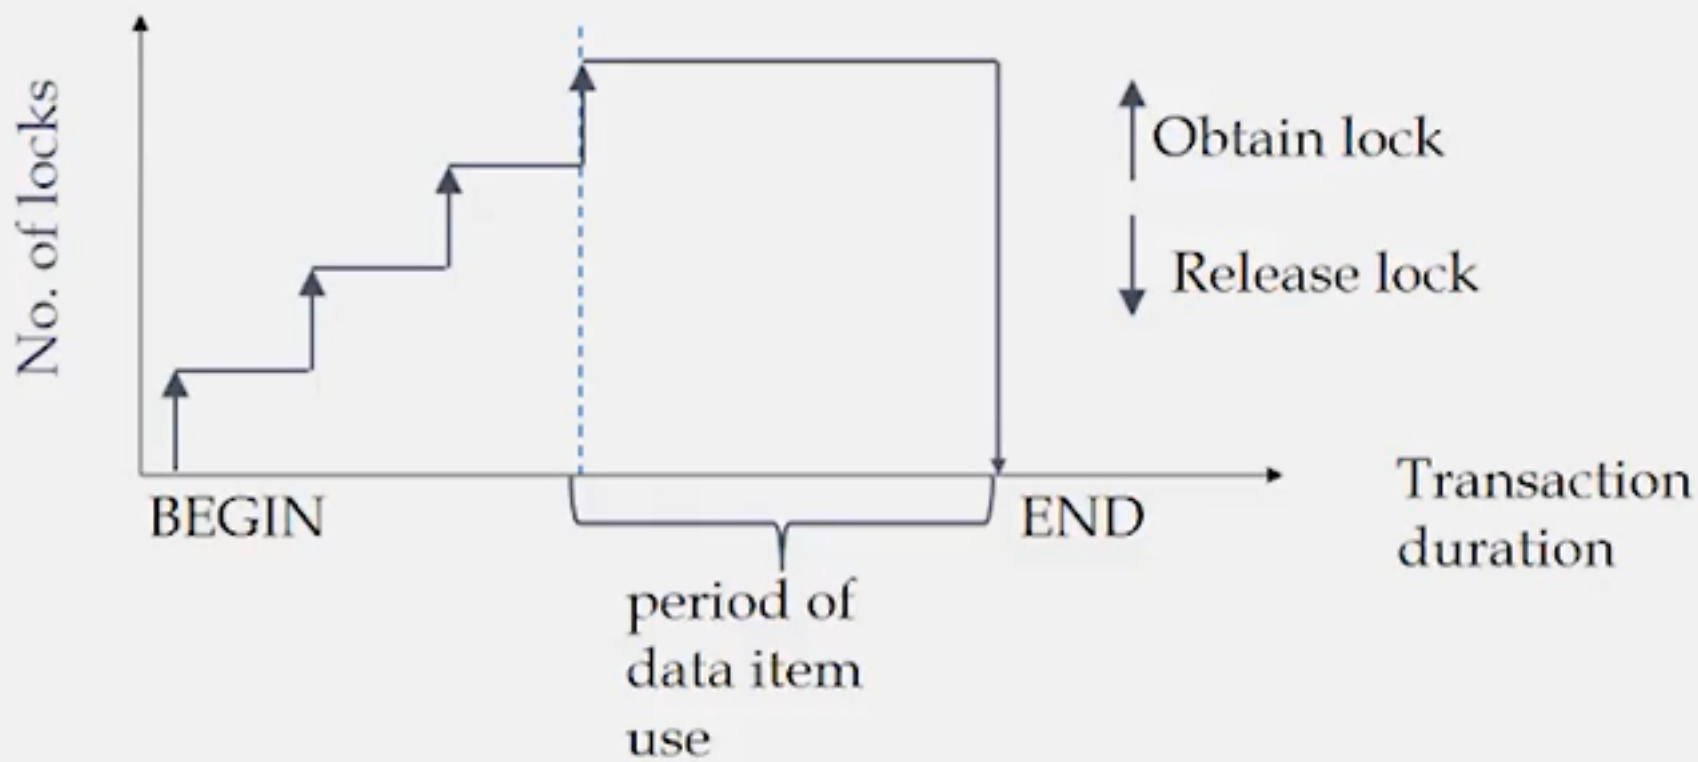
- concurrent changes invisible \Rightarrow serializable

DURABILITY

- committed updates persist

Strict 2PL


Hold locks until the end.



Deadlocks

- Deadlock: Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
 - Deadlock prevention ✓
 - Deadlock detection

Faults to Failures




Types of Failures

- Transaction failures
 - Transaction aborts (unilaterally or due to deadlock)
 - Avg. 3% of transactions abort abnormally
- System failures
 - Failure of processor, main memory, power supply, ...
 - Main memory contents are lost, but secondary storage contents are safe
 - Partial vs. total failure
- Media failures
 - Failure of secondary storage devices such that the stored data is lost
 - Head crash/controller failure (?)
- Communication failures
 - Lost/undeliverable messages
 - Network partitioning

Local Recovery Management – Architecture

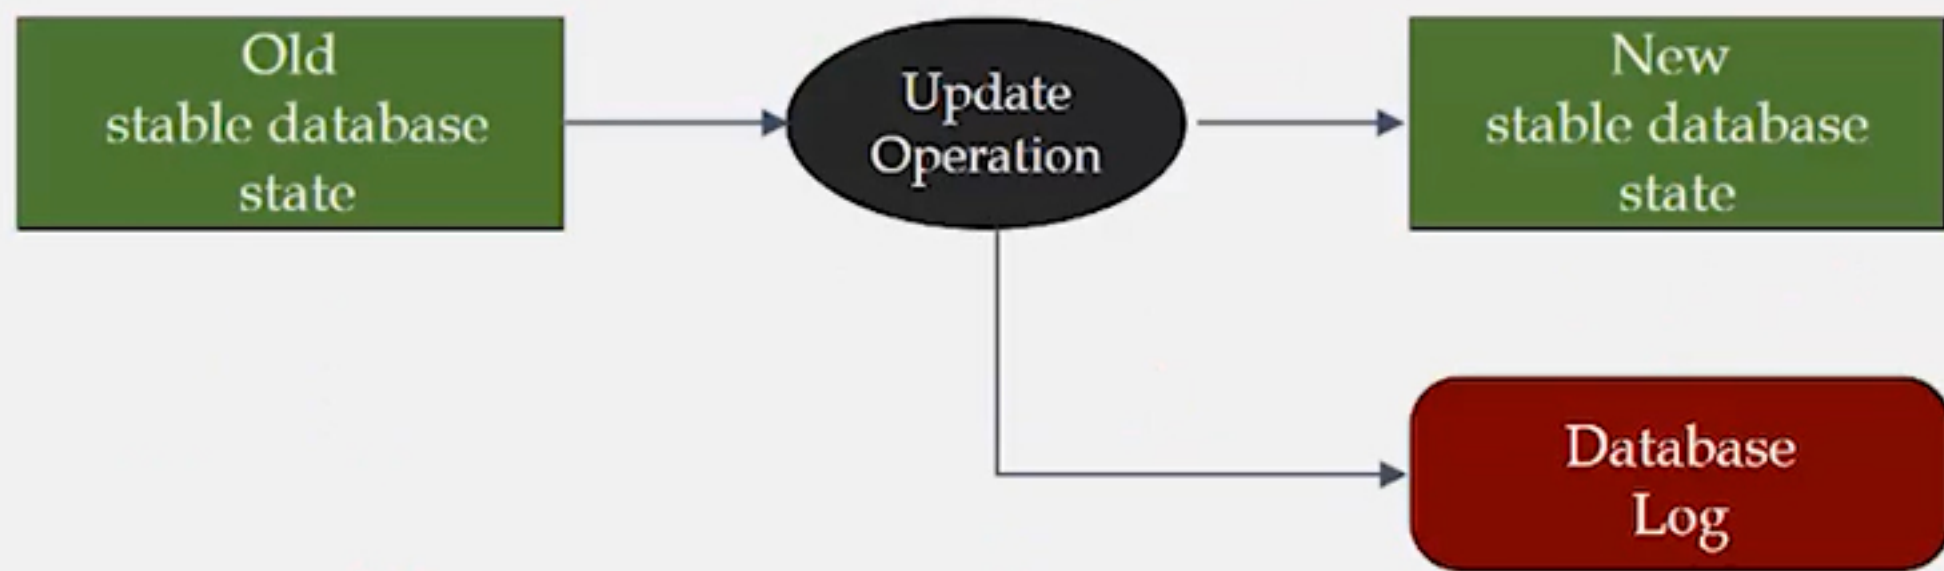
- **Volatile storage**
 - Consists of the main memory of the computer system (RAM).
- **Stable/Persistent storage**
 - Resilient to failures and loses its contents only in the presence of media failures (e.g., head crashes on disks).
 - Implemented via a combination of hardware (non-volatile storage) and software (stable-write, stable-read, clean-up) components.



In-Place Update Recovery Information

Database Log

Every action of a transaction must not only perform the action, but must also write a *log* record to an append-only file.



The Log

- The following actions are recorded in the log:
 - *Ti writes an object*: the old value and the new value.
 - Log record must go to disk before the changed page!
 - *Ti commits/aborts*: a log record indicating this action.
- Log records are chained together by Xact id, so it's easy to undo a specific Xact.
- All log related activities are handled transparently by the DBMS.

Logging

The log contains information used by the recovery process to restore the consistency of a system. This information may include


- transaction identifier
- type of operation (action)
- items accessed by the transaction to perform the action
- old value (state) of item (before image)
- new value (state) of item (after image)

...

Why Logging?

Upon recovery:


- all of T_1 's effects should be reflected in the database (REDO if necessary due to a failure)
- none of T_2 's effects should be reflected in the database (UNDO if necessary)



Write-Ahead Log Protocol


- Notice:
 - If a system crashes before a transaction is committed, then all the operations must be undone. Only need the before images (*undo portion* of the log).
 - Once a transaction is committed, some of its actions might have to be redone. Need the after images (*redo portion* of the log).
- WAL protocol :
 - ❶ Before a stable database is updated, the undo portion of the log should be written to the stable log
 - ❷ When a transaction commits, the redo portion of the log must be written to stable log prior to the updating of the stable database.

REDO Protocol



- REDO'ing an action means performing it again.
- The REDO operation uses the log information and performs the action that might have been done before, or not done due to failures.
- The REDO operation generates the new image.

UNDO Protocol



- UNDO'ing an action means to restore the object to its before image.
- The UNDO operation uses the log information and restores the old value of the object.

Summary

- Concurrency control and recovery are among the most important functions provided by a DBMS.
- Users need not worry about concurrency.
 - System automatically inserts lock/unlock requests and schedules actions of different Xacts in such a way as to ensure that the resulting execution is equivalent to executing the Xacts one after the other in some order.
- Write-ahead logging (WAL) is used to undo the actions of aborted transactions and to restore the system to a consistent state after a crash.
 - *Consistent state*: Only the effects of committed Xacts seen.